

# linear20

November 24, 2025

## 1 Applied Linear Regression: Auto Dataset Analysis

This notebook follows the assignment brief to analyze the `Auto` dataset. We will perform: 1. **Simple Linear Regression:** `mpg` as a function of `horsepower`. 2. **Multiple Linear Regression:** `mpg` as a function of all other relevant predictors.

We will load the data, fit the models, interpret the results, and produce diagnostic plots.

### 1.1 Setup: Import Libraries and Load Data

First, we import all the necessary Python libraries for data manipulation, statistical modeling, and plotting.

```
[1]: # --- Imports ---

# For data handling
import pandas as pd
import numpy as np

# For plotting
import matplotlib.pyplot as plt
import seaborn as sns

# For statistical models
import statsmodels.api as sm
import statsmodels.formula.api as smf
from statsmodels.stats.anova import anova_lm

# For diagnostics
import scipy.stats as stats
from statsmodels.graphics.regressionplots import influence_plot

# Set plot style
sns.set_theme(style="whitegrid")
# Show plots inline in Jupyter
%matplotlib inline
```

Now, we load the `Auto.csv` dataset from the provided URL.

**Important Data Cleaning Note:** The `horsepower` column in this dataset contains '?' for

missing values. We must: 1. Read the CSV and tell `pandas` to treat '?' as a Not-a-Number (NaN). 2. Drop any rows that have NaN values to ensure our regression models run correctly. 3. We will also drop the `name` column, as it's a unique identifier and not a useful predictor.

[4]: # --- Load Data ---

```
# FIX: This is the correct, stable URL for the raw CSV file
data_url = "https://www.statlearning.com/s/Auto.csv"

# Read the CSV, marking '?' as missing values
df = pd.read_csv(data_url, na_values='?')

# --- Clean Data ---

# Print original shape
print(f"Original shape: {df.shape}")

# Drop rows with any missing values (especially in 'horsepower')
df_clean = df.dropna()

# Print new shape
print(f"Shape after dropping NA: {df_clean.shape}")

# Drop the 'name' column as it's not a predictor
df_clean = df_clean.drop(columns=['name'])

# Display the first 5 rows and data types to confirm
print("\n--- Data Head ---")
print(df_clean.head())
print("\n--- Data Types ---")
print(df_clean.info())
```

Original shape: (397, 9)

Shape after dropping NA: (392, 9)

--- Data Head ---

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	\
0	18.0	8	307.0	130.0	3504	12.0	70	
1	15.0	8	350.0	165.0	3693	11.5	70	
2	18.0	8	318.0	150.0	3436	11.0	70	
3	16.0	8	304.0	150.0	3433	12.0	70	
4	17.0	8	302.0	140.0	3449	10.5	70	

origin

0	1
1	1
2	1
3	1

```

4      1

--- Data Types ---
<class 'pandas.core.frame.DataFrame'>
Index: 392 entries, 0 to 396
Data columns (total 8 columns):
 #   Column       Non-Null Count  Dtype  
---  --  
 0   mpg          392 non-null    float64 
 1   cylinders    392 non-null    int64  
 2   displacement 392 non-null    float64 
 3   horsepower   392 non-null    float64 
 4   weight        392 non-null    int64  
 5   acceleration 392 non-null    float64 
 6   year          392 non-null    int64  
 7   origin        392 non-null    int64  
dtypes: float64(4), int64(4)
memory usage: 27.6 KB
None

```

## 1.2 Task A: Simple Linear Regression (`mpg ~ horsepower`)

This task involves using simple linear regression with `mpg` as the response and `horsepower` as the predictor.

### 1.2.1 Task A(a): Fit and Summarize the Model

We use `statsmodels.formula.api.ols()` (Ordinary Least Squares) to fit the model. The formula '`mpg ~ horsepower`' automatically includes an intercept.

We will then: 1. Print the model summary. 2. Calculate the predicted `mpg` for `horsepower = 98`. 3. Get the 95% confidence and prediction intervals for that value.

```
[6]: # --- Task A(a): Fit and Summarize ---

# (i, ii, iii) Fit the model and print the summary
# We use df_clean which has no missing values
slr_model = smf.ols('mpg ~ horsepower', data=df_clean).fit()
print(slr_model.summary())

# (iv) Get prediction for horsepower = 98
# Create a new DataFrame for the value we want to predict
new_hp = pd.DataFrame({'horsepower': [98]})

# Get prediction
pred = slr_model.predict(new_hp)
print(f"\nPredicted mpg for horsepower=98: {pred.iloc[0]:.4f}")

# Get 95% confidence and prediction intervals
```

```

pred_summary = slr_model.get_prediction(new_hp).summary_frame(alpha=0.05)

print("\n--- 95% Intervals for horsepower=98 ---")
print(pred_summary)

```

OLS Regression Results

---

Dep. Variable:	mpg	R-squared:	0.606
Model:	OLS	Adj. R-squared:	0.605
Method:	Least Squares	F-statistic:	599.7
Date:	Mon, 24 Nov 2025	Prob (F-statistic):	7.03e-81
Time:	13:00:38	Log-Likelihood:	-1178.7
No. Observations:	392	AIC:	2361.
Df Residuals:	390	BIC:	2369.
Df Model:	1		
Covariance Type:	nonrobust		

---

	coef	std err	t	P> t	[0.025	0.975]
Intercept	39.9359	0.717	55.660	0.000	38.525	41.347
horsepower	-0.1578	0.006	-24.489	0.000	-0.171	-0.145

---

Omnibus:	16.432	Durbin-Watson:	0.920
Prob(Omnibus):	0.000	Jarque-Bera (JB):	17.305
Skew:	0.492	Prob(JB):	0.000175
Kurtosis:	3.299	Cond. No.	322.

---

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Predicted mpg for horsepower=98: 24.4671

```

--- 95% Intervals for horsepower=98 ---
      mean   mean_se  mean_ci_lower  mean_ci_upper  obs_ci_lower  \
0  24.467077  0.251262      23.973079      24.961075      14.809396

      obs_ci_upper
0      34.124758

```

**Your Interpretation for A(a):** (*This is where you'll write your comments based on the output above*)

- i. Is there a relationship?
  - Hint: Look at the  $P>|t|$  (p-value) for horsepower. Is it very small (e.g.,  $< 0.05$ )?
- ii. How strong is the relationship?
  - Hint: Look at the R-squared value. This represents the percentage of variance in mpg

*explained by horsepower.*

- iii. Is the relationship positive or negative?

- Hint: Look at the `coef` for `horsepower`. A negative sign means as `horsepower` goes up, `mpg` goes down.

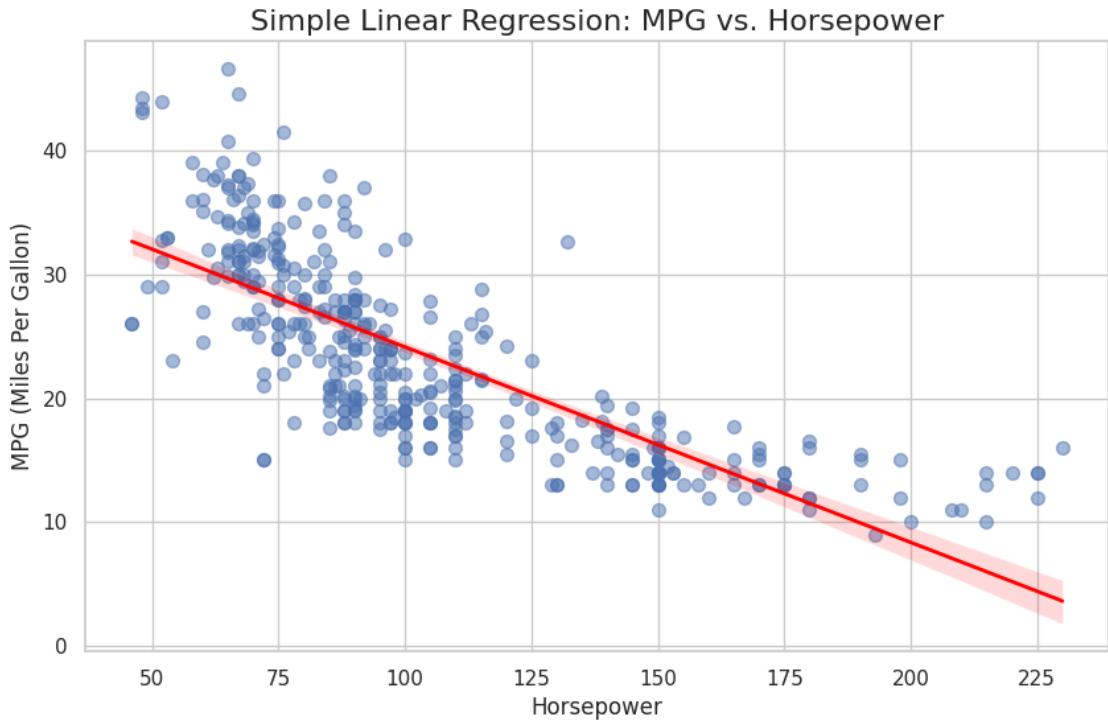
- iv. What is the predicted mpg?

- The output gives the exact value. The **confidence interval** (e.g., `mean_ci_lower`, `mean_ci_upper`) is for the average\* mpg for a car with 98 hp. The **prediction interval** (e.g., `obs_ci_lower`, `obs_ci_upper`) is for a *single specific* car with 98 hp, so it's wider.\*

### 1.2.2 Task A(b): Plot Response and Predictor

We'll use Seaborn's `regplot` to easily create a scatter plot and overlay the least squares regression line.

```
[7]: # --- Task A(b): Plot ---  
  
plt.figure(figsize=(10, 6))  
sns.regplot(  
    x='horsepower',  
    y='mpg',  
    data=df_clean,  
    line_kws={'color': 'red', 'linewidth': 2},  # Style for the regression line  
    scatter_kws={'alpha': 0.5, 's': 50}          # Style for the scatter points  
)  
plt.title('Simple Linear Regression: MPG vs. Horsepower', fontsize=16)  
plt.xlabel('Horsepower', fontsize=12)  
plt.ylabel('MPG (Miles Per Gallon)', fontsize=12)  
plt.show()
```



### 1.2.3 Task A(c): Diagnostic Plots

We need to check the assumptions of the linear model. We'll look at two key plots: 1. **Residuals vs. Fitted Values:** To check for non-linearity (i.e., patterns like a U-shape) and non-constant variance (heteroscedasticity). 2. **Normal Q-Q Plot:** To check if the residuals are normally distributed.

```
[8]: # --- Task A(c): Diagnostics ---

# Get fitted values and residuals from the model
fitted_vals = slr_model.fittedvalues
residuals = slr_model.resid

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 6))

# 1. Residuals vs. Fitted Plot
sns.residplot(
    x=fitted_vals,
    y=residuals,
    lowess=True, # Adds a smooth line to see trends
    ax=ax1,
    line_kws={'color': 'red', 'lw': 2},
    scatter_kws={'alpha': 0.5}
)
```

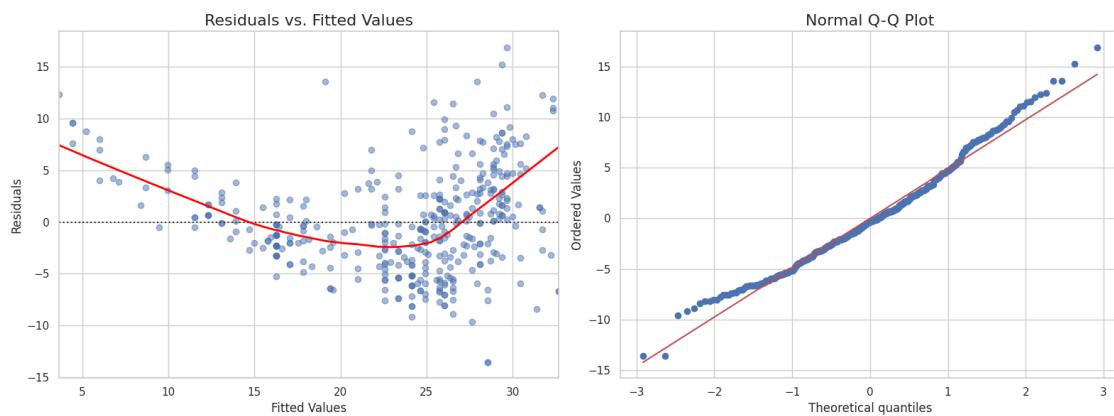
```

ax1.set_title('Residuals vs. Fitted Values', fontsize=16)
ax1.set_xlabel('Fitted Values', fontsize=12)
ax1.set_ylabel('Residuals', fontsize=12)

# 2. Normal Q-Q Plot
stats.probplot(residuals, dist="norm", plot=ax2)
ax2.set_title('Normal Q-Q Plot', fontsize=16)

plt.tight_layout()
plt.show()

```



**Your Commentary for A(c):** (*This is where you'll write your comments based on the plots above*)

- Hint: Look at the **Residuals vs. Fitted** plot. Do you see a clear curve or U-shape? This would suggest the relationship is **not linear**, and a simple straight line isn't the best fit. Does the spread of the points get wider as the fitted values increase? That would be heteroscedasticity.
- Hint: Look at the **Q-Q Plot**. Do the blue dots fall nicely along the red line? If they curve off at the ends, it suggests the residuals are not perfectly normally distributed.

### 1.3 Task B: Multiple Linear Regression

Now we'll build a model to predict mpg using all other variables as predictors.

#### 1.3.1 Task B(a): Scatterplot Matrix

A scatterplot matrix (or pairplot) shows the relationship between every pair of variables.

[10]: # --- Task B(a): Scatterplot Matrix ---

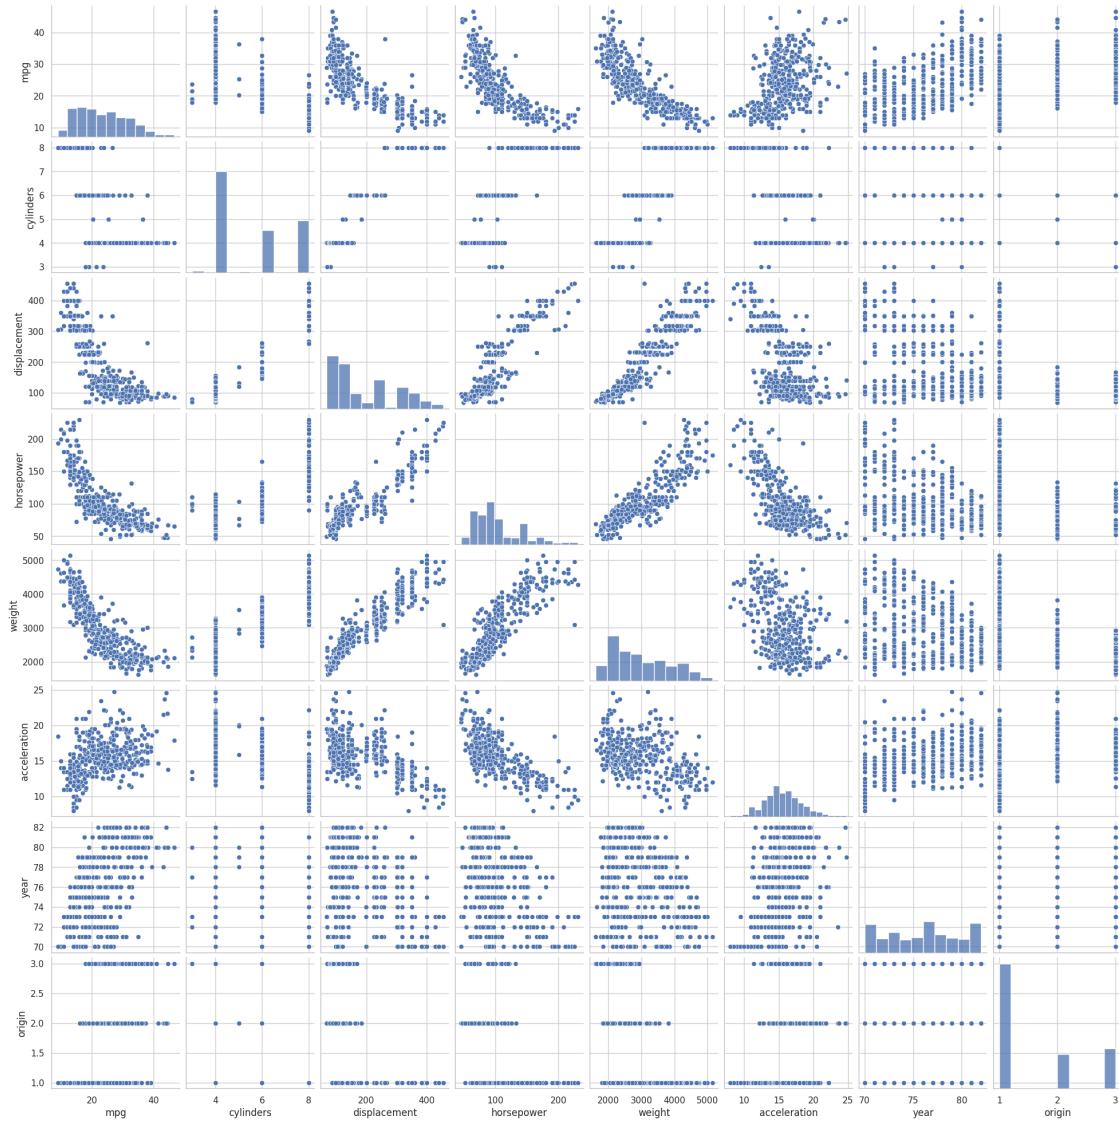
```

print("Generating pairplot... (this may take a moment)")
# We use df_clean (which doesn't have the 'name' column)
sns.pairplot(df_clean)

```

```
plt.show()
```

Generating pairplot... (this may take a moment)



### 1.3.2 Task B(b): Correlation Matrix

Let's compute the exact correlation coefficients and visualize them with a heatmap.

```
[11]: # --- Task B(b): Correlation Matrix ---
```

```
# Compute the correlation matrix
corr_matrix = df_clean.corr()
```

```

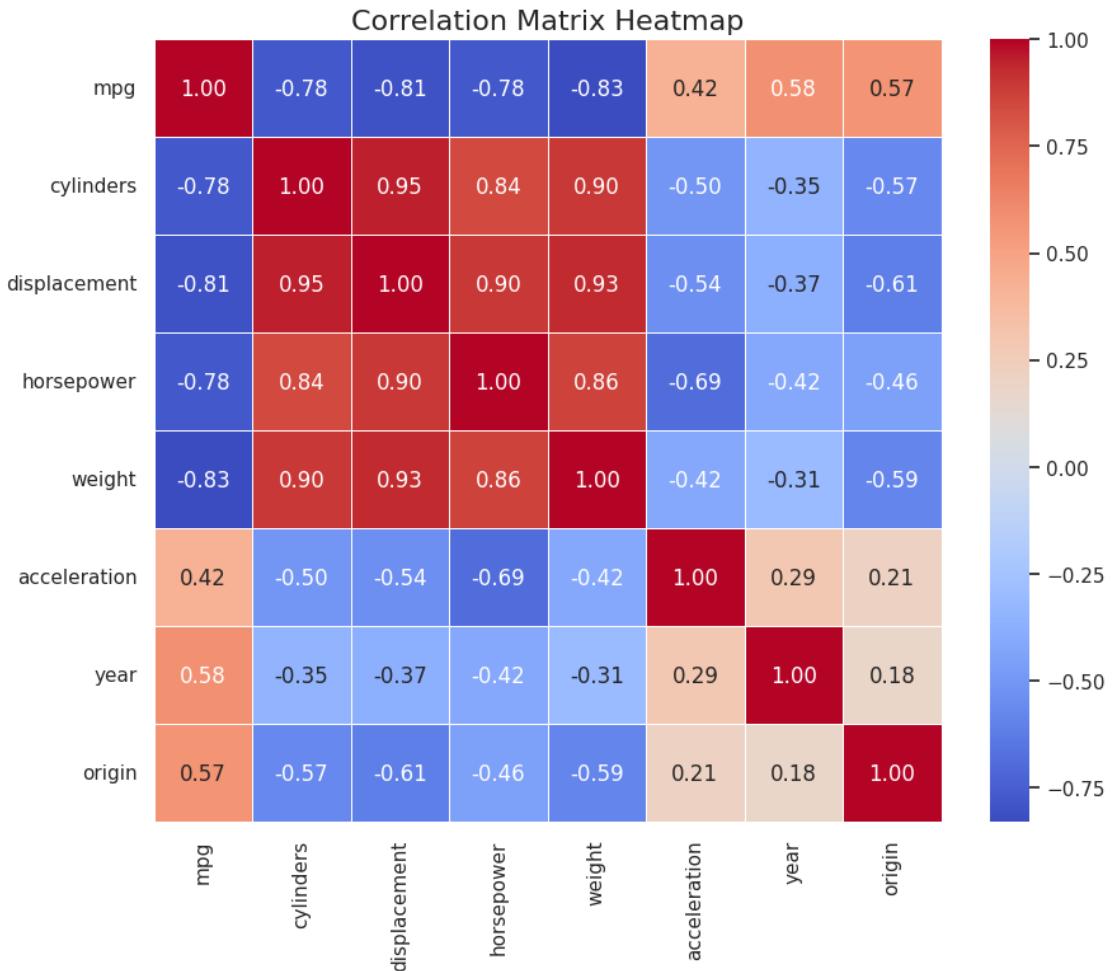
print("--- Correlation Matrix ---")
print(corr_matrix)

# Plot the heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(
    corr_matrix,
    annot=True,      # Show the numbers in the cells
    cmap='coolwarm', # Color scheme
    fmt='.2f',        # Format numbers to 2 decimal places
    linewidths=0.5
)
plt.title('Correlation Matrix Heatmap', fontsize=16)
plt.show()

```

--- Correlation Matrix ---

	mpg	cylinders	displacement	horsepower	weight	\
mpg	1.000000	-0.777618	-0.805127	-0.778427	-0.832244	
cylinders	-0.777618	1.000000	0.950823	0.842983	0.897527	
displacement	-0.805127	0.950823	1.000000	0.897257	0.932994	
horsepower	-0.778427	0.842983	0.897257	1.000000	0.864538	
weight	-0.832244	0.897527	0.932994	0.864538	1.000000	
acceleration	0.423329	-0.504683	-0.543800	-0.689196	-0.416839	
year	0.580541	-0.345647	-0.369855	-0.416361	-0.309120	
origin	0.565209	-0.568932	-0.614535	-0.455171	-0.585005	
	acceleration	year	origin			
mpg	0.423329	0.580541	0.565209			
cylinders	-0.504683	-0.345647	-0.568932			
displacement	-0.543800	-0.369855	-0.614535			
horsepower	-0.689196	-0.416361	-0.455171			
weight	-0.416839	-0.309120	-0.585005			
acceleration	1.000000	0.290316	0.212746			
year	0.290316	1.000000	0.181528			
origin	0.212746	0.181528	1.000000			



### 1.3.3 Task B(c): Fit and Summarize the Model

We fit the multiple linear regression model. The formula `mpg ~ .` is a shortcut that means “use all other columns in the DataFrame as predictors.” We’ll also run an ANOVA test.

```
[12]: # --- Task B(c): Fit and Summarize ---
```

```
# The formula 'mpg ~ .' uses all other columns as predictors
# Since df_clean only contains our desired columns, this is safe.
all_predictors = ' + '.join(df_clean.columns.drop('mpg'))
formula = f'mpg ~ {all_predictors}'

print(f"Using formula: {formula}\n")

# Fit the model
mlr_model = smf.ols(formula, data=df_clean).fit()
```

```

# Print the summary
print(mlr_model.summary())

# (i) Run ANOVA test to check overall model significance
anova_table = anova_lm(mlr_model, typ=2)
print("\n--- ANOVA Table ---")
print(anova_table)

```

Using formula: mpg ~ cylinders + displacement + horsepower + weight + acceleration + year + origin

#### OLS Regression Results

Dep. Variable:	mpg	R-squared:	0.821			
Model:	OLS	Adj. R-squared:	0.818			
Method:	Least Squares	F-statistic:	252.4			
Date:	Mon, 24 Nov 2025	Prob (F-statistic):	2.04e-139			
Time:	13:04:58	Log-Likelihood:	-1023.5			
No. Observations:	392	AIC:	2063.			
Df Residuals:	384	BIC:	2095.			
Df Model:	7					
Covariance Type:	nonrobust					
<hr/>						
	coef	std err	t	P> t	[0.025	0.975]
<hr/>						
Intercept	-17.2184	4.644	-3.707	0.000	-26.350	-8.087
cylinders	-0.4934	0.323	-1.526	0.128	-1.129	0.142
displacement	0.0199	0.008	2.647	0.008	0.005	0.035
horsepower	-0.0170	0.014	-1.230	0.220	-0.044	0.010
weight	-0.0065	0.001	-9.929	0.000	-0.008	-0.005
acceleration	0.0806	0.099	0.815	0.415	-0.114	0.275
year	0.7508	0.051	14.729	0.000	0.651	0.851
origin	1.4261	0.278	5.127	0.000	0.879	1.973
<hr/>						
Omnibus:	31.906	Durbin-Watson:			1.309	
Prob(Omnibus):	0.000	Jarque-Bera (JB):			53.100	
Skew:	0.529	Prob(JB):			2.95e-12	
Kurtosis:	4.460	Cond. No.			8.59e+04	
<hr/>						

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 8.59e+04. This might indicate that there are strong multicollinearity or other numerical problems.

--- ANOVA Table ---

sum_sq	df	F	PR(>F)
--------	----	---	--------

cylinders	25.791491	1.0	2.329125	1.277965e-01
displacement	77.612668	1.0	7.008884	8.444649e-03
horsepower	16.739754	1.0	1.511699	2.196328e-01
weight	1091.631693	1.0	98.580813	7.874953e-21
acceleration	7.358417	1.0	0.664509	4.154780e-01
year	2402.249906	1.0	216.937408	3.055983e-39
origin	291.134494	1.0	26.291171	4.665681e-07
Residual	4252.212530	384.0	NaN	NaN

**Your Interpretation for B(c):** (*This is where you'll write your comments based on the output above*)

- i. Is there a relationship?
  - Hint: Look at the Prob (F-statistic) in the main summary or the PR(>F) for the model in the ANOVA table. A very small value means at least one predictor is related to mpg.
- ii. Which predictors are significant?
  - Hint: In the main summary table, look at the P>/t/ column for each predictor. Which ones are < 0.05? displacement, weight, year, and origin are likely candidates.
- iii. What does the year coefficient suggest?
  - Hint: The coefficient (e.g., ~0.75) means that for each additional model year, the mpg is expected to **increase** by 0.75, holding all other variables constant.

### 1.3.4 Task B(d): Diagnostic Plots

We repeat the diagnostic plots for the multiple regression model. We also add an **Influence Plot** to find:

- \* **Outliers:** Points with high residuals (far from 0 on the y-axis).
- \* **High-Leverage Points:** Points that have unusual predictor values (far from 0 on the x-axis). These points can have a strong pull on the regression line.

```
[14]: # --- Task B(d): Diagnostics ---

# Get fitted values and residuals
mlr_fitted_vals = mlr_model.fittedvalues
mlr_residuals = mlr_model.resid

# 1. Residuals vs. Fitted
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 6))
sns.residplot(
    x=mlr_fitted_vals,
    y=mlr_residuals,
    lowess=True,
    ax=ax1,
    line_kws={'color': 'red', 'lw': 2},
    scatter_kws={'alpha': 0.5}
)
ax1.set_title('Residuals vs. Fitted Values (MLR)', fontsize=16)
ax1.set_xlabel('Fitted Values', fontsize=12)
```

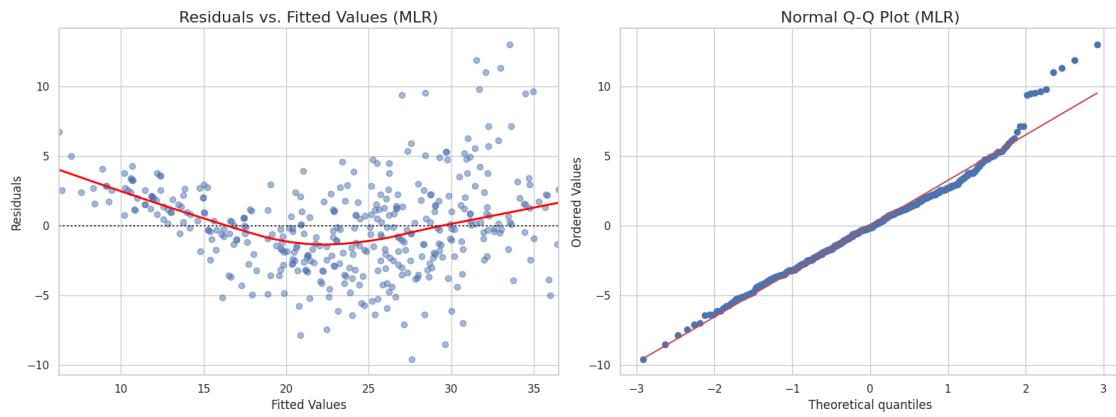
```

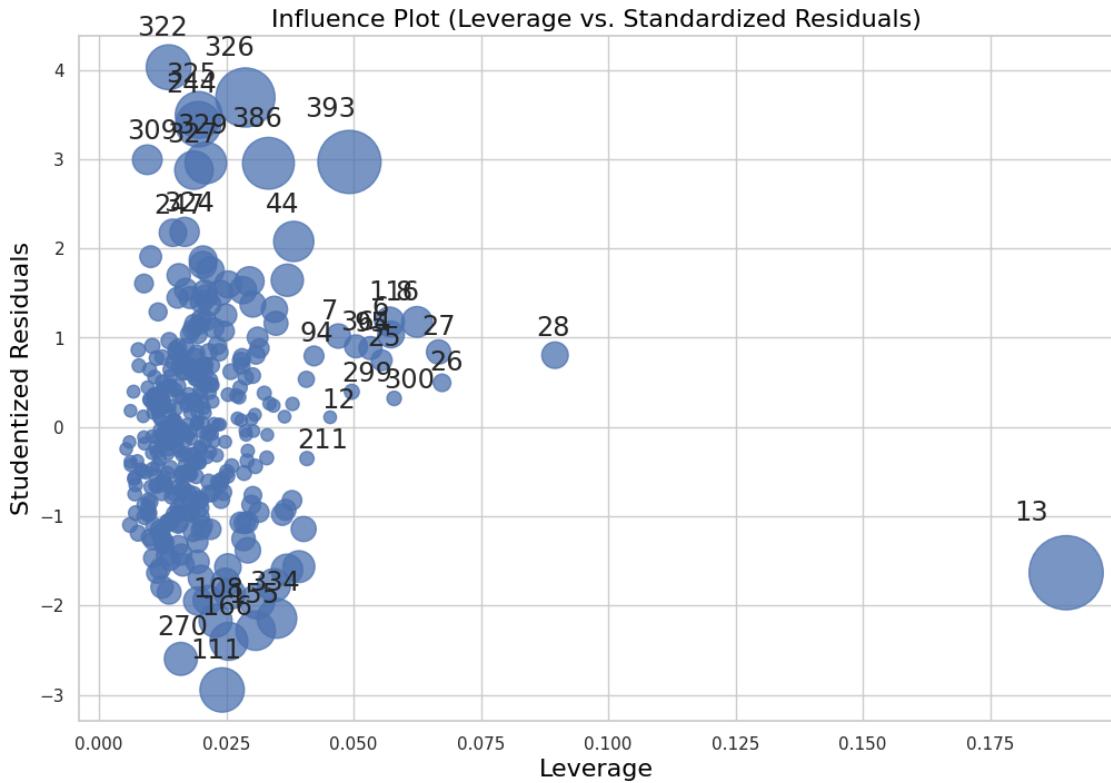
ax1.set_ylabel('Residuals', fontsize=12)

# 2. Normal Q-Q Plot
stats.probplot(mlr_residuals, dist="norm", plot=ax2)
ax2.set_title('Normal Q-Q Plot (MLR)', fontsize=16)
plt.tight_layout()
plt.show()

# 3. Influence Plot (Leverage vs. Residuals)
fig, ax = plt.subplots(figsize=(12, 8))
influence_plot(mlr_model, ax=ax, criterion="cooks")
plt.title('Influence Plot (Leverage vs. Standardized Residuals)', fontsize=16)
plt.show()

```





**Your Commentary for B(d):** (*This is where you'll write your comments based on the plots above*)

- Hint: Does the **Residuals vs. Fitted** plot look better or worse than the simple model? Is the curve less pronounced?
- Hint: In the **Influence Plot**, look for points with high leverage (far right on x-axis) or large residuals (far up/down on y-axis). The size of the bubble indicates Cook's distance (a measure of overall influence). Are there any points that are both high-leverage and high-residual? The point labeled 13 or 14 often stands out as high leverage.

### 1.3.5 Task B(e): Fit Models with Interactions

Let's see if the effect of one variable depends on another. We can test an interaction term using the \* or : syntax. \* a \* b includes a, b, and the interaction a:b. \* a : b includes *only* the interaction.

Let's test `horsepower * weight`.

[15]: # --- Task B(e): Interactions ---

```
# We'll use all predictors, but add an interaction between horsepower and weight
# We use '*' which includes horsepower, weight, AND their interaction ↴
        ↴ (horsepower:weight)
```

```

interaction_formula = 'mpg ~ cylinders + displacement + horsepower * weight +_
    ↴acceleration + year + origin'

interaction_model = smf.ols(interaction_formula, data=df_clean).fit()
print(interaction_model.summary())

```

OLS Regression Results

---

Dep. Variable:	mpg	R-squared:	0.862
Model:	OLS	Adj. R-squared:	0.859
Method:	Least Squares	F-statistic:	298.6
Date:	Mon, 24 Nov 2025	Prob (F-statistic):	1.88e-159
Time:	13:06:34	Log-Likelihood:	-973.24
No. Observations:	392	AIC:	1964.
Df Residuals:	383	BIC:	2000.
Df Model:	8		
Covariance Type:	nonrobust		

---

=====

	coef	std err	t	P> t	[0.025
0.975]					
-----	-----	-----	-----	-----	-----
Intercept	2.8757	4.511	0.638	0.524	-5.993
11.744					
cylinders	-0.0296	0.288	-0.103	0.918	-0.596
0.537					
displacement	0.0059	0.007	0.881	0.379	-0.007
0.019					
horsepower	-0.2313	0.024	-9.791	0.000	-0.278
-0.185					
weight	-0.0112	0.001	-15.393	0.000	-0.013
-0.010					
horsepower:weight	5.529e-05	5.23e-06	10.577	0.000	4.5e-05
6.56e-05					
acceleration	-0.0902	0.089	-1.019	0.309	-0.264
0.084					
year	0.7695	0.045	17.124	0.000	0.681
0.858					
origin	0.8344	0.251	3.320	0.001	0.340
1.329					

---

Omnibus:	40.936	Durbin-Watson:	1.474
Prob(Omnibus):	0.000	Jarque-Bera (JB):	73.199
Skew:	0.629	Prob(JB):	1.27e-16
Kurtosis:	4.703	Cond. No.	1.23e+07

---

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.23e+07. This might indicate that there are strong multicollinearity or other numerical problems.

**Your Commentary for B(e):** (*This is where you'll write your comments based on the output above*)

- Hint: Look at the summary table for the new interaction term, `horsepower:weight`. Is its  $P>/t/$  value significant (e.g.,  $< 0.05$ )? If yes, this suggests the effect of `horsepower` on `mpg` depends on the weight of the car (and vice-versa).

### 1.3.6 Task B(f): Variable Transformations

The residual plots in Task A showed a non-linear pattern. This suggests we might get a better fit by transforming our predictors. Let's try `log(X)`, `sqrt(X)`, and `X^2`.

We use `I()` in the formula to perform calculations like `I(horsepower**2)`.

[17]: # --- Task B(f): Transformations ---

```
# We'll build a few different models to compare
# We use np.log() and I() directly in the formulas

# 1. Log transformation on horsepower
log_model = smf.ols('mpg ~ np.log(horsepower) + weight + year + origin', □
    ↵data=df_clean).fit()
print("--- Model with log(horsepower) ---")
print(log_model.summary())

# 2. Quadratic transformation (polynomial) on horsepower
# We include both horsepower and horsepower^2
poly_model = smf.ols('mpg ~ horsepower + I(horsepower**2) + weight + year +' □
    ↵'origin', data=df_clean).fit()
print("\n--- Model with horsepower^2 ---")
print(poly_model.summary())
```

--- Model with log(horsepower) ---

OLS Regression Results

Dep. Variable:	mpg	R-squared:	0.828
Model:	OLS	Adj. R-squared:	0.826
Method:	Least Squares	F-statistic:	464.5
Date:	Mon, 24 Nov 2025	Prob (F-statistic):	2.91e-146
Time:	13:07:45	Log-Likelihood:	-1016.6
No. Observations:	392	AIC:	2043.
Df Residuals:	387	BIC:	2063.
Df Model:	4		

```

Covariance Type: nonrobust
=====
=====

            coef      std err          t      P>|t|      [0.025
0.975]
-----
Intercept      4.3697      6.095      0.717      0.474     -7.614
16.354
np.log(horsepower) -4.9700      1.040     -4.780      0.000     -7.014
-2.926
weight        -0.0043      0.000     -9.741      0.000     -0.005
-0.003
year           0.6925      0.049     14.154      0.000      0.596
0.789
origin         1.2467      0.253      4.929      0.000      0.749
1.744
=====
Omnibus:      28.598   Durbin-Watson:      1.314
Prob(Omnibus): 0.000   Jarque-Bera (JB): 42.111
Skew:          0.525   Prob(JB):       7.17e-10
Kurtosis:      4.215   Cond. No.      1.16e+05
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.16e+05. This might indicate that there are strong multicollinearity or other numerical problems.

--- Model with horsepower^2 ---

OLS Regression Results

```

Dep. Variable:      mpg      R-squared:      0.851
Model:             OLS      Adj. R-squared: 0.849
Method:            Least Squares      F-statistic: 439.5
Date:      Mon, 24 Nov 2025      Prob (F-statistic): 7.11e-157
Time:          13:07:45      Log-Likelihood: -988.57
No. Observations: 392      AIC:           1989.
Df Residuals:    386      BIC:           2013.
Df Model:          5
Covariance Type: nonrobust
=====
=====

            coef      std err          t      P>|t|      [0.025
0.975]
-----

```

Intercept	-6.6457	3.915	-1.698	0.090	-14.343
1.052					
horsepower	-0.2441	0.027	-9.099	0.000	-0.297
-0.191					
I(horsepower ** 2)	0.0008	9.13e-05	9.170	0.000	0.001
0.001					
weight	-0.0044	0.000	-10.426	0.000	-0.005
-0.004					
year	0.7456	0.046	16.145	0.000	0.655
0.836					
origin	1.0465	0.238	4.405	0.000	0.579
1.514					
<hr/>					
Omnibus:	21.819	Durbin-Watson:			1.500
Prob(Omnibus):	0.000	Jarque-Bera (JB):			32.447
Skew:	0.414	Prob(JB):			9.00e-08
Kurtosis:	4.140	Cond. No.			4.10e+05
<hr/>					

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 4.1e+05. This might indicate that there are strong multicollinearity or other numerical problems.

**Your Commentary for B(f):** (*This is where you'll write your comments based on the output above*)

- Hint: Compare the *R-squared* and *Adj. R-squared* of these new models to the original multiple regression model (*mlr\_model*). Are they higher? A higher *Adj. R-squared* indicates a better model fit.
- Look at the *p*-values for the new terms. Is *np.log(horsepower)* significant? Are both *horsepower* and *I(horsepower\*\*2)* significant in the polynomial model? If the *I(horsepower\*\*2)* term is significant, it strongly confirms that the relationship between *horsepower* and *mpg* is non-linear (quadratic).
- For a complete analysis, you would also re-run the diagnostic plots (Residuals vs. Fitted) for your best\* new model to confirm that the non-linear pattern is gone.\*

## 1.4 Reflective Summary

(*This is where you write your final summary for your blog post*)

- Example: “In this analysis, we explored the *Auto* dataset to predict fuel efficiency (*mpg*). A simple linear regression showed a strong, negative relationship between *horsepower* and *mpg*, but diagnostic plots revealed a clear non-linear pattern. A multiple regression model identified *weight*, *year*, and *origin* as the most significant predictors. By testing transformations, we found that a model including a quadratic term (*I(horsepower\*\*2)*) or a log term (*np.log(horsepower)*) provided a better fit, as indicated by a higher Adjusted *R-squared* and more significant terms. This confirms that the relationship between *horsepower* and fuel effi-

*ciency is not a simple straight line. The diagnostic plots for the transformed model showed... which suggests... A key takeaway is the importance of...*