

# Untitled

December 5, 2025

```
[7]: import pandas as pd
import numpy as np
from sklearn.ensemble import GradientBoostingRegressor, RandomForestRegressor, VotingRegressor
from sklearn.model_selection import KFold

# 1. Load Data
train_df = pd.read_csv('train.csv')
test_df = pd.read_csv('test.csv')

# -----
# SIMPLE PREPROCESSING (Back to Basics)
# -----
def process_data(df):
    # Fix Horsepower
    df['horsepower'] = pd.to_numeric(df['horsepower'], errors='coerce')
    df['horsepower'] = df['horsepower'].fillna(df['horsepower'].median())

    # Physics: Log Weight is the most important feature
    df['log_weight'] = np.log(df['weight'])

    # Physics: Power-to-Weight
    df['power_to_weight'] = df['horsepower'] / df['weight']

    return df

train_df = process_data(train_df)
test_df = process_data(test_df)

# -----
# FEATURES
# -----
# Trees handle 'origin' and 'cylinders' as numbers just fine.
# We don't need to change them to categories.
features = ['displacement', 'horsepower', 'log_weight', 'acceleration', 'year', 'origin', 'cylinders', 'power_to_weight']
```

```

X_train_full = train_df[features]
y_train_full = train_df['mpg']
X_test = test_df[features]

# Safety Check: Fill gaps
X_train_full = X_train_full.fillna(X_train_full.median())
X_test = X_test.fillna(X_train_full.median())

# -----
# DEFINE THE STABLE MODELS
# -----

def get_model():
    # Model 1: Gradient Boosting (The Smart One)
    gbr = GradientBoostingRegressor(
        n_estimators=500,
        learning_rate=0.02, # Slow and steady learning prevents "exploding"
        scores
        max_depth=3,
        random_state=42
    )

    # Model 2: Random Forest (The Stable One)
    # Random Forest is very hard to confuse. It averages 200 random decision
    trees.
    rf = RandomForestRegressor(
        n_estimators=200,
        max_depth=10,
        random_state=42
    )

    return VotingRegressor([('gbr', gbr), ('rf', rf)])

# -----
# 10-FOLD BAGGING (To smooth out errors)
# -----

folds = 10
kf = KFold(n_splits=folds, shuffle=True, random_state=42)
test_predictions_sum = np.zeros(len(X_test))

print(f"Starting {folds}-Fold Training (Recovery Mode)...")

for fold, (train_index, val_index) in enumerate(kf.split(X_train_full, y_train_full)):

    X_fold = X_train_full.iloc[train_index]
    y_fold = y_train_full.iloc[train_index]

```

```

# Log Transform Target (Still the best trick for this dataset)
y_fold_log = np.log(y_fold)

model = get_model()
model.fit(X_fold, y_fold_log)

log_pred = model.predict(X_test)
test_predictions_sum += np.exp(log_pred)

print(f" - Fold {fold+1} complete.")

# Average
final_predictions = test_predictions_sum / folds

# -----
# SAVE
# -----
submission = pd.DataFrame({
    'ID': test_df['ID'],
    'mpg': final_predictions
})

submission.to_csv('submission.csv', index=False)

print("\n Success! New 'submission.csv' generated.")
print(submission.head())

```

Starting 10-Fold Training (Recovery Mode)...

- Fold 1 complete.
- Fold 2 complete.
- Fold 3 complete.
- Fold 4 complete.
- Fold 5 complete.
- Fold 6 complete.
- Fold 7 complete.
- Fold 8 complete.
- Fold 9 complete.
- Fold 10 complete.

Success! New 'submission.csv' generated.

		ID	mpg
0	70_chevrolet chevelle malibu_alpha_3505	16.487742	
1	71_buick skylark 320_bravo_3697	14.606418	
2	70_plymouth satellite_charlie_3421	16.242916	
3	68_amc rebel sst_delta_3418	16.164524	
4	70_ford torino_echo_3444	16.651944	